


Bias and reasoning performance



Bias error averages available normalized bias metrics; reasoning averages ARC-Challenge, ARC-Easy, and Balanced COPA. Lower bias and higher reasoning accuracy are better. Existing-baseline reasoning uses Llama-70B baseline utility runs when available.